

Source-to-source automatic differentiation using PSyclone

A prototype implementation

AIRSEA team meeting
September 5, 2023

Julien REMY (M1 student - Martin's trainee)

julien.remy@grenoble-inp.org

Source code: <https://github.com/JulienRemy/PSyclone>


Documentation: <https://psyclone-autodiff.readthedocs.io/> (not complete yet)



Automatic differentiation

Program

Suppose we have a program computing return values of variables y_j from values of arguments x_i through intermediate values v_k , where each value is obtained from its direct predecessors through *elemental* operations ($+$, \times , $/$, \exp , etc.).



Notations

Relation: $i \prec j$ if v_j depends on v_i , eg. $1 \prec 5$.

Predecessors: $u_j := (v_i)_{i \prec j}$ eg. $u_5 = \begin{pmatrix} v_1 \\ v_2 \end{pmatrix}$

Operation: $\varphi_j : u_j \mapsto v_j$ eg. $v_5 = \varphi_5(u_5)$

AD: Obtain $\frac{\partial y_j}{\partial x_i}(x_1, \dots, x_n)$ automatically, by differentiating operations $\varphi_k : u_k \mapsto v_k$.

Forward mode Tangent

Derivatives

For **one independent variable** z chosen among the **argument** variables, say $z = x_1$,

let
$$\dot{v}_i = \frac{\partial v_i}{\partial z}$$

$$\dot{v}_{k+2} = \frac{\partial v_{k+2}}{\partial v_{k+1}} \underbrace{\left(\frac{\partial v_{k+1}}{\partial v_k} \dot{v}_k \right)}_{\dot{v}_{k+1}}$$

The chain rule gives

Easy to implement.

$$\dot{v}_j = \sum_{i \prec j} \frac{\partial \varphi_j}{\partial v_i}(u_j) \dot{v}_i$$

Initialize $\dot{z} = 1$,

and $\dot{v}_k = 0, \forall v_k \neq z$,

get $\left(\frac{\partial y_j}{\partial z}(x_1, \dots, x_n) \right)_j$

Issue

We obtain $J((x_i)_i)(0 \dots \dot{z} \dots 0)^T$: all the derivatives $(\dot{y}_j)_j$ wrt a **single** $z \in (x_i)_i$.

Interesting if $\#\{x_i\}_i \ll \#\{y_j\}_j$ but usually there are many arguments, few results.

Reverse mode Adjoint

Adjoint

For **one dependent variable** z chosen among the **return** variables, say $z = y_1$, denote

$$\bar{v}_i = \frac{\partial z}{\partial v_i} \qquad \underbrace{\left(\bar{v}_k \frac{\partial v_k}{\partial v_{k-1}} \right)}_{\bar{v}_{k-1}} \frac{\partial v_{k-1}}{\partial v_{k-2}} = \bar{v}_{k-2}$$

The chain rule gives

In practice, increment.

$$\bar{v}_i = \sum_{j \succ i} \bar{v}_j \frac{\partial \varphi_j}{\partial v_i}(\overset{?}{u}_j)$$

Initialize $\bar{z} = 1$,

and $\bar{v}_k = 0, \forall v_k \neq z$,

get $\left(\frac{\partial z}{\partial x_i}(x_1, \dots, x_n) \right)_i$


Remark

We obtain the adjoints $(\bar{x}_i)_i$ of **all arguments** x_i for a single differentiated $z \in (y_j)_j$.


Compute $\nabla^T z(x_1, \dots, x_n) = (0 \dots \bar{z} \dots 0) J(x_1, \dots, x_n)$ in one evaluation!

Reverse mode AD

A simple math example, with non linearities



Initialize with $\forall i, \bar{x}_i = 0, \quad \forall k, \bar{v}_k = 0$ and choose $(\bar{y}_1 = 1, \bar{y}_2 = 0)$ **or** $(\bar{y}_1 = 0, \bar{y}_2 = 1)$



What about programs?

First issue

Overwrites: taping

Recording routine

```
b = sin(a)
tape(i) = a
a = x + 1 ! overwrite a
c = cos(a)
```

Returning routine

```
! here a == x + 1
a_adj = a_adj + c_adj * (-sin(a));      c_adj = 0.0
x_adj = x_adj + a_adj * 1;              a_adj = 0.0
a = tape(i)
a_adj = a_adj + b_adj * cos(a);          b_adj = 0.0
```

Taping

Record and restore the values of overwritten variables to a "tape".
Usually a LIFO stack. But then we can't parallelize. Instead, I'm using (static) arrays.
Other possibility: recompute values, optional checkpointing.

Taping results of operations

Can also be used for the results of composed operations. Not implemented yet.
PSyclone: Operation nodes have no datatypes for now, especially shapes (for vectors).

What about programs?

Second issue

Iterative assignments

Consider the following Fortran code and its naive and **wrong** reverse-mode AD:

$a = 2 * a + x$
! now reverse

$a_adj = a_adj + a_adj * 2$! wrong! $\bar{a} \not\rightarrow 3\bar{a}$
 $x_adj = x_adj + a_adj * 1$! wrong too!

Solution: deal with the LHS adjoint last

$a = 2 * a + x$
! now reverse

$x_adj = x_adj + a_adj * 1$! correct
 $a_adj = a_adj * 2$! correct

Reversal schedules 1/2

Nested calls

Suppose we want to differentiate the subroutine foo, which calls bar and qux.

```
subroutine foo(x, y)
  call bar(x, y)
  call qux(x, y)
end subroutine foo
```

Reversal schedules


Let us call and denote :

- advancing \square routine the original,
- recording \triangleright routine the one recording values to the tape,
- returning \triangleleft routine the one computing the adjoints,
- reversing $\triangleright\triangleleft$ routine the one combining the two preceding. Call it to differentiate.

Reversal schedules 2/2


Split reversal

Larger tape in memory, but computations are only run once.



Joint reversal

Smaller tape but computations are repeated.



"Link" reversal

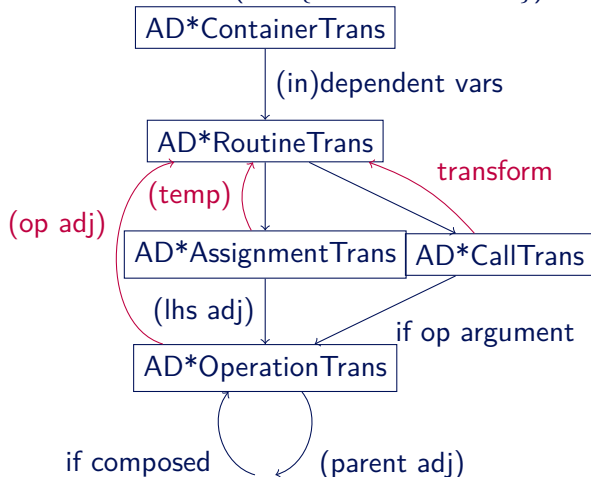
Specify whether to use a split or joint reversal for each (caller → called) pair.

AD prototype module in PSyclone - psyclone.autodiff

Forward- and reverse-mode:

- scalars only,
- subroutines only (no functions),
- calls to subroutines,
- reverse: joint, split and "link" reversals,
- adjoints and derivatives of fixed type and precision,
- quite naive implementation:
 - operations are always recomputed, differentiating $v = 1/u$ or $v = \exp(u)$ doesn't use the postvalue of v .
- simplification and substitution as a post-processing step.

Transformations: ($*$ \in {Forward, Reverse})



Test suite and planned features

Existing tests

- Unit tests for (most) elemental transformations.
- Transformations outputs tested by comparing to Tapenade (to be automated)
 - correct up to associativity/commutativity floating point errors in both modes.

Planned features

- Vectors and matrices,
- Activity (\leftrightarrow dependence DAG) and TBR (to-be-recorded) analyses,
- Loops: need to be run backward in reverse-mode:
 - except if trivial: record the loop indices, run-time length of the tape
- Control flow: record the boolean values of conditions? Tag the branches?
- Functions? Programs?

Thank you for attending!
Questions? Remarks? Ideas?

The logo for Inria, featuring the word "Inria" in a stylized, red, cursive script.The logo for PSyclone, featuring the word "PSyclone" in a green, sans-serif font, followed by a stylized graphic of three black, curved lines with green dots, resembling a DNA helix or a stylized 'S'.